

1. What assumptions does linear regression make?

Answer:

1.
Linearity: The relationship between each predictor and the response is linear.
2.
Independence: Observations (and residuals) are independent.
3.
Homoscedasticity: Residuals have constant variance across predictor values.
4.
Normality of Errors: Residuals (errors) are normally distributed (important for inference).
5.
No or Low Multicollinearity: Predictors are not highly correlated with each other.

2. How do you interpret the coefficients?

Answer:

- A coefficient represents the expected change in the response variable for a one-unit increase in predictor , holding other predictors constant.
- **Sign:** Positive indicates a direct relationship; negative indicates an inverse relationship.
- **Magnitude:** Larger absolute values imply stronger effects (e.g., means each unit increase in increases the prediction by 2.5 units).

3. What is R^2 score and its significance?

Answer:

- R^2 (coefficient of determination) measures the proportion of variance in the response explained by the model.
- Range: 0 to 1; values closer to 1 indicate a better fit.
- **Adjusted R^2** is preferred in multiple regression to account for the number of predictors.

4. When would you prefer MSE over MAE?

Answer:

- MSE (Mean Squared Error) penalizes larger errors more heavily (squaring amplifies outliers).
- MSE is differentiable, making it suitable for analytic solutions and gradient-based optimization.
- Use MAE (Mean Absolute Error) when robustness to outliers is more important.

5. How do you detect multicollinearity?

Answer:

- Variance Inflation Factor (VIF): A $VIF > 5-10$ indicates problematic collinearity.

- **Correlation Matrix:** Look for predictor pairs with high absolute correlation ().
- **Condition Number** of the design matrix: Values > 30 signal ill-conditioning.

6. What is the difference between simple and multiple regression?

Answer:

- **Simple Regression:** One predictor ().
- **Multiple Regression:** Two or more predictors ().
- Multiple regression models more complex relationships and controls for confounders.

7. Can linear regression be used for classification?

Answer:

- Not directly, as it predicts continuous outcomes.
- Thresholding its predictions can produce a classifier (e.g., $>0.5 \rightarrow$ class 1), but this may yield values outside $[0,1]$ and lacks probabilistic interpretation.
- Logistic regression is preferred for classification tasks.

8. What happens if you violate regression assumptions?

Answer:

- Biased/Inefficient Estimates: E.g., heteroscedasticity inflates variance of estimates.
-

Invalid Inference: P-values and confidence intervals become unreliable.

-
- Poor Generalization:** Misspecified models may not predict well on new data.
-
- Remedies:** Transformations, robust methods, regularization, or revisiting model specification.